


Modern C++ Design

Chapter 10: Visitor

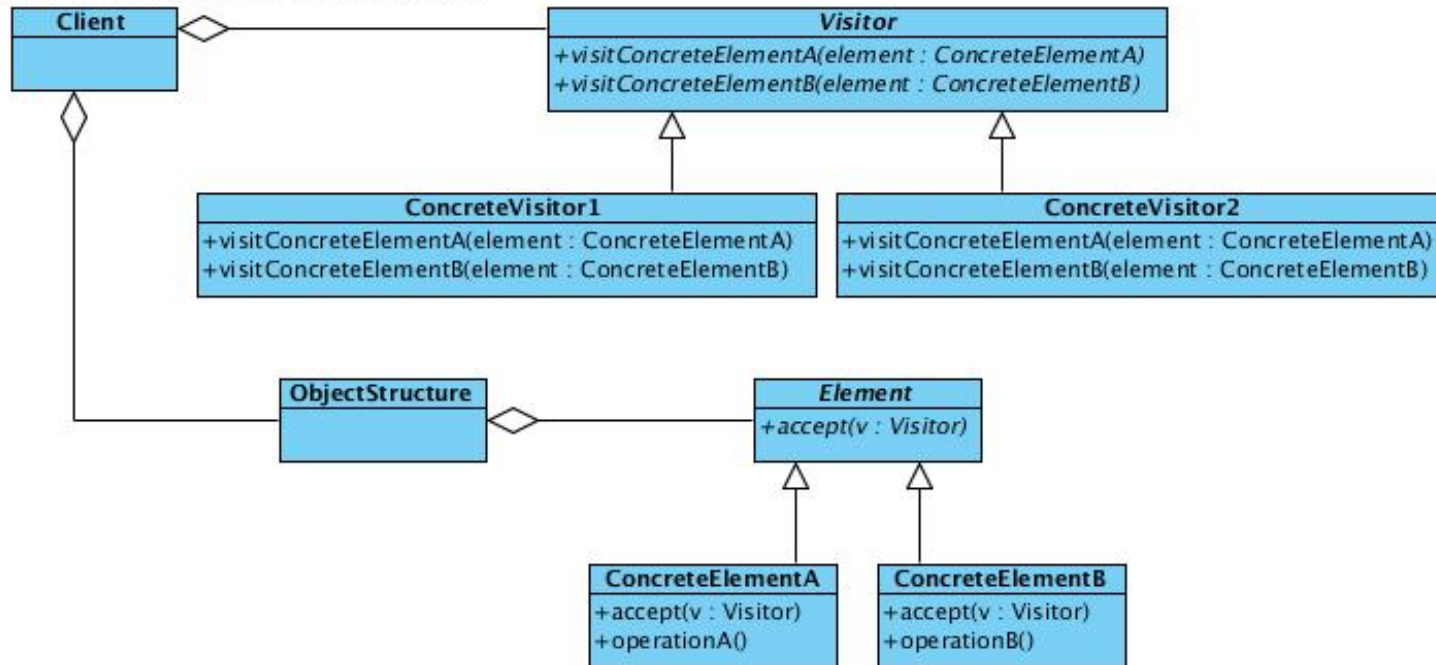


Introduction

Visitor gives a lot of flexibility at the expense of disabling features that designers take for granted.

- it enables addition of behaviour (new **ConcreteVisitor**) without adding code to where the behaviour is to be applied or recompiling other classes in the Visitor hierarchy or the clients.
- a leaf class cannot be added to the **Element** hierarchy without recompiling the Visitor hierarchy and all its clients.

Visitor works well with stable hierarchies and with heavy processing needs



```

returnType ConcreteVisitor1::visit(
    ConcreteElementA element) {
    element.operationA();
}
returnType ConcreteElementA::accept(Visitor v) {
    v.visit(this);
}

```

Goal of the Chapter

‘Craft’ a dependable version of Visitor, leaving as little burden on the application programmer as possible.

After reading the chapter, you should:

- understand how the Visitor works
- know when to and when not to use the Visitor pattern
- understand the drawbacks of the GoF implementation and know how to overcome these
- understand a generic implementation of the Visitor and be able to apply it to problems requiring a Visitor pattern.

Visitor Basics

In general adding new functionality to a hierarchy can be done in two ways:

- By adding a new class
- By adding a new virtual member function

Adding a new class is easy - no need to recompile existing classes. Good code reuse.

Adding a new virtual function is difficult - to make use of polymorphic properties you may need to add a virtual function in the base class and possibly other classes. This requires a recompilation of the 'world' to take place.

This is in contrast to what the visitor allows....

Consider a Document Editor example

- Document elements, such as paragraphs, vector drawings, bitmaps etc. are represented as classes and derive from a common root, **DocElement**
- A document is a structured collection of pointers to **DocElements**.
- You wish to iterate through the document structure and perform operations such as spellchecking, reformatting, gathering of document statistics...
- This type of operation requires code to be added, not modified. The addition of code in one place will ease code maintenance.

- Consider the following Document Statistics class:

```
class DocStats {  
    ...  
public:  
    void AddChars(unsigned int charsToAdd) {  
        chars_ += charsToAdd;  
    }  
    // Similarly for AddWords, AddImages ...  
  
    // Display the statistics to the user in  
    // a dialog box  
    void Display();  
private:  
    unsigned int chars_, nonBlankChars_,  
                words_, images_;  
};
```

Using a classic OO approach to gather the stats will result in a virtual function in `DocElement` that deals with gathering stats

```
class DocElement {  
    ...  
    // This member function helps with the  
    // "Statistics" feature  
    virtual void UpdateStats(DocStats& statistics)=0;  
};
```


Each concrete document element defines the function in its own way.

```
class Paragraph :public DocElement {
    ...
    void UpdateStats(DocStats& statistics) {
        statistics.AddChars(number of chars
                               in the paragraph);
        statistics.AddWords(number of words
                               in the paragraph);
    }
    ... }

class RasterBitMap :public DocElement {
    ...
    void UpdateStats(DocStats& statistics) {
        statistics.AddImages(1);
    }
    ... }
```

The ‘driver’ function will be defined as follows:

```
void Document::DisplayStatistics() {  
    DocStats statistics;  
    foreach (DocElement in the document as element) {  
        element->UpdateStats(statistics);  
    }  
    statistics.Display();  
}
```

Disadvantages of this approach:

- **DocElement** and all classes deriving from it must have access to the **DocStats** definition. Every change to **DocStats** will result in the **DocElement** hierarchy needing to be recompiled.
- The operations for gathering the stats are spread throughout the **UpdateStats** implementations. Maintenance of the stats feature requires that multiple classes/files need to be considered.
- Adding additional operations that relate to stats gathering requires a virtual function to be added to **DocElement**. This implies adding the function to all classes that derive from **DocElement** as well.

Breaking the dependency of **DocElement** on **DocStats** is necessary.

1. Move all operations into **DocStats**.

- **DocStats** must then figure out what to do with each concrete type.
- **DocStats** therefore requires a member function `void UpdateStats(DocElement&)`.
- Document now iterates through its elements and calls **UpdateStats** for each of them.
- **DocStats** is now invisible to **DocElement**, however **DocStats** now depends on each **DocElement** that it needs to process.
- **UpdateStats** now relies on a *type switch* - call a polymorphic object on its concrete type and perform different operations with it depending on that concrete type.

```

void DocStats::UpdateStats(DocElement& elem) {
    if (Paragraph* p =
        dynamic_cast<Paragraph*>(&elem)) {
        chars_ += p->NumChars();
        words_ += p->NumWords();
    } else
        if (dynamic_cast<RasterBitmap*>(&elem) {
            ++images_;
        }
    else ...
}

```

- This code is hard to understand, hard to extend, and hard to maintain.
- The order in which the *dynamic casts* are done are of utmost importance - Refer to Chapter 8

2. Make use of the visitor pattern

- you require new functions to act as virtual without adding a new virtual function for each operation
- make use of what is referred to as a *unique bouncing virtual function* in the **DocElement** hierarchy. This function ‘teleports’ work to another hierarchy. The **DocElement** hierarchy is referred to as the *visited* hierarchy and the operations belong to the new *visitor* hierarchy.
- “Each implementation of the bouncing virtual function calls a *different* function in the visitor hierarchy.”, thereby selecting the visited types. “The functions in the visitor hierarchy called by the bouncing function are virtual.”, thereby selecting the operations.

The DocElementVisitor class

```
class DocElementVisitor {  
    public:  
        virtual void VisitParagraph(Paragraph&)=0;  
        virtual void VisitRasterBitmap(RasterBitmap&)=0;  
        ...  
};
```

Add the bouncing virtual function (**Accept**) to the DocElement hierarchy. It takes a DocElementVisitor& and invokes the appropriate VisitXxxx function.

```
class DocElement {  
    public:  
        virtual void Accept(DocElementVisitor&) = 0;  
        ...  
};
```

```
void Paragraph::Accept(DocElementVisitor& v) {  
    v.VisitParagraph(*this);  
}
```

```
void RasterBitmap::Accept(DocElementVisitor& v) {  
    v.VisitRasterBitmap(*this);  
}
```

DocStats now changes

```
class DocStats : public DocElementVisitor {  
public:  
    virtual void VisitParagraph(Paragraph& par) {  
        chars_ += par.NumChars();  
        words_ += par.NumWords();  
    }  
}
```



```

    virtual void VisitRasterBitmap(RasterBitmap&) {
        ++images;
    }
    ...
};

```

The ‘driver’ function changes as follows:

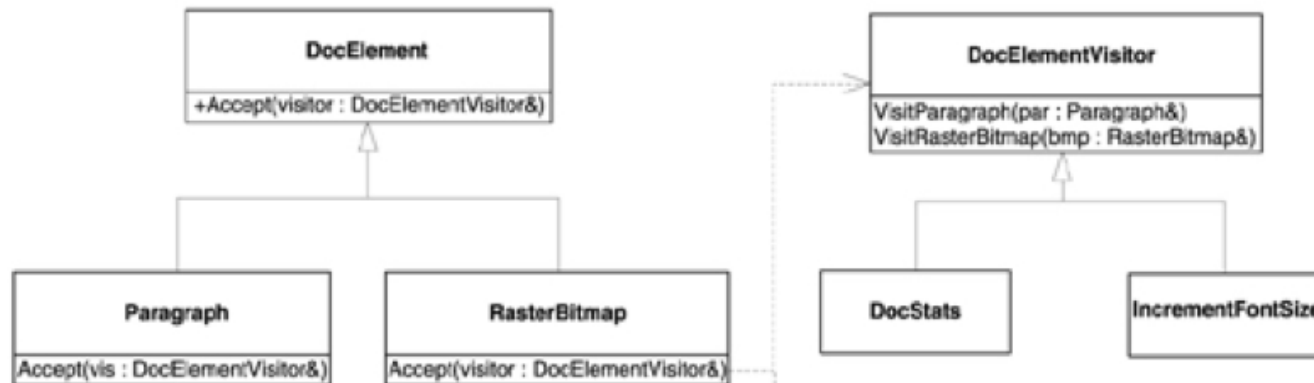
```

void Document::DisplayStatistics() {
    DocStats statistics;
    foreach (DocElement in the document as element) {
        // element->UpdateStats(statistics);
        element->Accept(statistics)
    }
    statistics.Display();
}

```

Adding a visitor, for example to increase font size is easy!

```
class IncrementFontSize : public DocElementVisitor{  
    virtual void VisitParagraph(Paragraph& par) {  
        par.SetFontSize(par.GetFontSize() + 1);  
    }  
    virtual void VisitRasterBitmap(RasterBitmap&){  
        // do nothing  
    }  
};
```



For each document element type, the `DocElementVisitor` defined a virtual function with the element type as part of the function name. This can be simplified by letting the compiler determine the overloaded function to call.

```
class DocElementVisitor {
public:
    // virtual void VisitParagraph(Paragraph&) = 0;
    virtual void Visit(Paragraph&) = 0;
    // virtual void VisitRasterBitmap(RasterBitmap&) = 0;
    virtual void Visit(RasterBitmap&) = 0;
    ...
};
```

The `Accept` functions for each of the document elements also simplifies.

```
void Paragraph::Accept(DocElementVisitor& v) {
    // v.VisitParagraph(*this);
    v.Visit(*this); }
}
```

```

void RasterBitmap::Accept(DocElementVisitor& v) {
    //v.VisitRasterBitmap(*this);
    v.Visit(*this); }

```

NOTE: the static type of `*this` in `Paragraph::Accept` is `Paragraph&`, and similarly for `RasterBitmap`. This means these functions are quite different and cannot be factored out.

What-if there is no matching Visit function for the type received and you do not want a compile-time error! Can provide a catch-all as the last resort.

```

class DocElementVisitor {
public:
    // virtual void VisitParagraph(Paragraph&) = 0;
    virtual void Visit(Paragraph&) = 0;
    // virtual void VisitRasterBitmap(RasterBitmap&) = 0;
    virtual void Visit(RasterBitmap&) = 0;

```

```
...  
virtual void Visit(DocElement&) = 0;  
};
```

The Acyclic Visitor

The previous example exhibits cyclic dependencies... `DocElement` needs `DocElementVisitor`, and `DocElementVisitor` needs all of `DocElement`'s hierarchy

Dividing the classes into files needs to be done with care

```
// File DocElementVisitor.h
class DocElement;
class Paragraph;
class RasterBitmap;
    ... forward declarations for all
        DocElement derivatives

class DocElementVisitor {
    public:
        virtual void Visit(Paragraph&) = 0;
        ...
};
```

```
// File DocElement.h  
class DocElementVisitor;  
  
class DocElement {  
    public:  
        virtual void Accept(DocElementVisitor&) = 0;  
        ...  
};
```


Adding new classes to the `DocElement` hierarchy becomes complex - something we are not meant to do as the premise was that this hierarchy remains stable. But it may be necessary at a stage to do so, for example adding the document element `VectorGraphic`.

- Add a new forward declaration for `VectorGraphic` in `DocElementVisitor.h`

- Add a new pure overload to `DocElementVisitor`

```
class DocElementVisistor {  
    ...  
    virtual void Visit(VectorGraphic&) = 0;  
};
```

- Every concrete visitor must implement `Visit(VectorGraphic&)` as a do-nothing as opposed to a pure-virtual function
- Implement `Accept` in the `VectorGraphic` class
- Recompile all the `DocElement` and `DocElementVisitor` hierarchies.

Removing cycles

- In 1996, Robert Martin developed a variation of the Visitor pattern to remove cycles.
- Develop a base class for the visitor hierarchy that only carries type information
- The visited hierarchy's **Accept** function accepts a reference to this base class visitor and applies a *dynamic cast* to detect the matching visitor
- A match results in a jump from the visited hierarchy to the visitor hierarchy

How is this implemented?

- Define the archetypical base class, the strawman

```
class DocElementVisitor {  
    public:  
        virtual ~DocElementVisitor() {}  
};
```

- The virtual destructor does 2 things: (i) gives `DocElementVisitor` RTTI capabilities; and (ii) it ensures correct polymorphic destruction of `DocElementVisitor` objects.
- `DocElement` class remains that same
- for each derived class of `DocElement` a visiting class is defined with one function `VisitXxxx`

```
class ParagraphVisitor {  
    public:  
        virtual void VisitParagraph(Paragraph&) = 0;  
};
```
- `Paragraph::Accept` changes to

```
void Paragraph::Accept(DocElementVisitor& v) {  
    if (ParagraphVisitor* p =  
        dynamic_cast<ParagraphVisitor*>(&v)) {  
        p->VisitParagraph(*this);  
    } else {
```

```

        // call a catch-all function
    }
}

```

- Other document element classes define similar implementations of **Accept**.
- A concrete visitor now derives from **DocElementVisitor** and all archetypical visitors of all classes that is of interest to it.

```

class DocStats :
    public DocElementVisitor, // required
    public ParagraphVisitor,
    public RasterBitmapVistor {
public:
    virtual void VisitParagraph(Paragraph& par)
        chars_ += par.NumChars();
        words_ += par.NumWords();
}

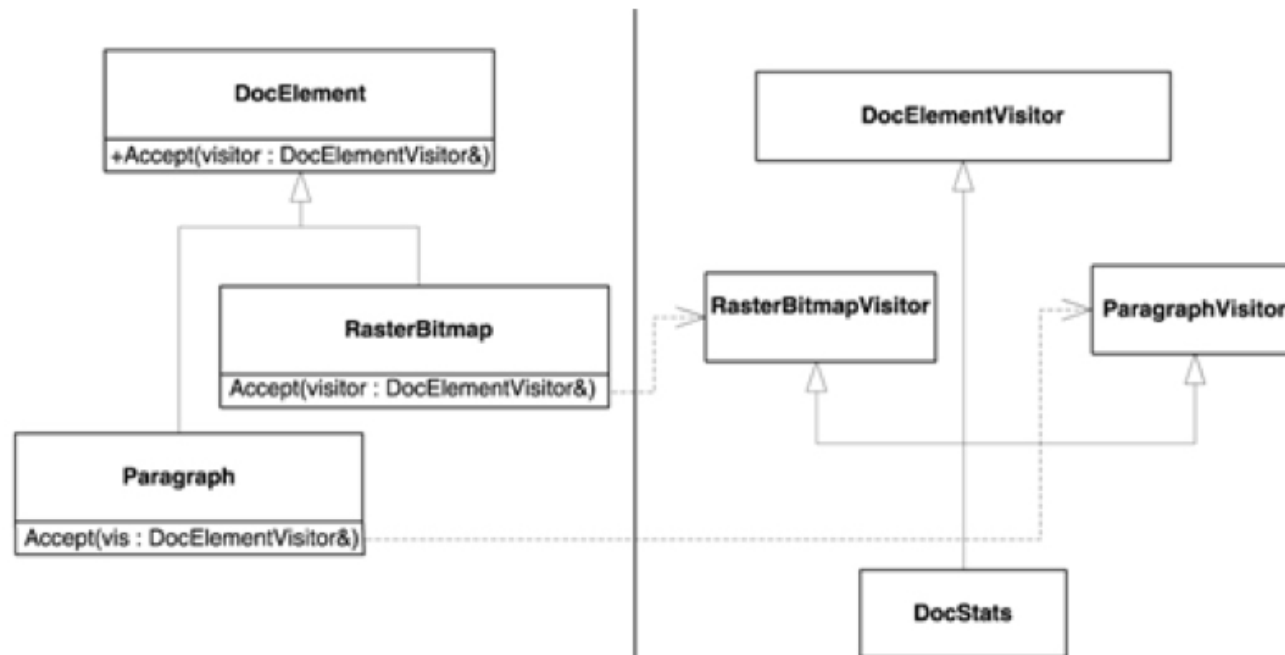
```

```

    virtual void VisitRasterBitmap(RasterBitmap&)
        ++images;
}
...
};

```

- The resulting class structure is given by



Cycles have been removed, but parallel hierarchies need to be maintained.

A Generic Implementation of Visitor

Divide the implementation into 2 units

- Visitable classes - the hierarchy to be visited (add operations to)
- Visitor classes - where the actual operations are implemented

Refer to Table 10.1 for Component names

Table 10.1. Component Names		
Name	Belongs To	Represents
BaseVisitable	Library	The root of all hierarchies that can be visited
Visitable	Library	A mix-in class that confers visitable properties to a class in the visitable hierarchy
DocElement	Application	A sample class—the root of the hierarchy we want to make visitable
Paragraph, RasterBitmap	Application	Two sample visitable concrete classes, derived from DocElement
Accept	Library and application	The member function that's called for visiting the visitable hierarchy
BaseVisitor	Library	The root of the visitor hierarchy
Visitor	Library	A mix-in class that confers visitor properties to a class in the visitor hierarchy
Statistics	Application	A sample visitor concrete class
Visit	Library and application	The member function that is called back by the visitable elements, for its visitors

Begin by defining the *Visitor* hierarchy - provide a base class and classes that define the **Visit** operation for the given type

```
class BaseVisitor {
    public:
        virtual ~BaseVisitor() {}
};

template <class T, typename R = void>
class Visitor {
    public:
        typedef R ReturnType;
        virtual ReturnType Visit(T&) = 0;
};
```

To visit the hierarchy you derive your visitor from `BaseVisitor` and all the `Visitor` instantiations as types you want to visit

```
class SomeVisitor :
    public BaseVisitor, // required
    public Visitor<Paragraph>,
    public Visitor<RasterBirmap> {
public:
    void Visit(Paragraph&);
    void Visit(RasterBitmap&);
};
```


Define the *Visitable* hierarchy

- Define the base for the hierarchy

```
template <typename R = void>
class BaseVisitable {
public:
    typedef R ReturnType;
    virtual ~BaseVisitable() {}
    virtual ReturnType Accept(BaseVisitor&) = 0;
};
```

- Need to find a way to implement **Accept** in the framework. This is done with macros - the client must insert **DEFINE_VISITABLE()** in each class of the visitable hierarchy. The decision is made in the name of *efficiency*

```
#define DEFINE_VISITABLE() \
    virtual Return Type Accept(BaseVisitor& guest) \
    { return AcceptImpl(*this, guest); }
```

- the macro defines **Accept** to forward to another function **AcceptImpl**
- **AcceptImpl** is parameterised with the type of ***this**
- **AcceptImpl** therefore gains access to the static type.
- class **BaseVisitable** changes to implement **AcceptImpl**

```

template <typename R = void>
class BaseVisitable {
public:
    typedef R ReturnType;
    virtual ~BaseVisitable() {}
    virtual ReturnType Accept(BaseVisitor&)=0;
protected:
    // give access only to the hierarchy
    template <class T>
    static ReturnType AcceptImpl(T& visited,
                                   BaseVisitor& guest) {
        // Apply the acyclic visitor
        if (Visitor<T>* p =
            dynamic_cast<Visitor<T>*>(&guest)) {
            return p->Visit(visited);
        }
    }
}

```

```
        return ReturnType();  
    } };
```

- The design of the *Visitor* and *Visitable* is given by:

```
// Visitor part  
class BaseVisitor  
{  
public:  
    virtual ~BaseVisitor() {}  
};  
template <class T, typename R = void>  
class Visitor  
{  
public:  
    typedef R ReturnType; // Available for clients  
    virtual ReturnType Visit(T&) = 0;  
};
```

```

// Visitable part
template <typename R = void>
class BaseVisitable
{
public:
    typedef R ReturnType;
    virtual ~BaseVisitable() {}
    virtual R Accept(BaseVisitor&) = 0;
protected:
template <class T>
static ReturnType AcceptImpl(T& visited,
                             BaseVisitor& guest) {
    // Apply the Acyclic Visitor
    if (Visitor<T>* p =
        dynamic_cast<Visitor<T>*>(&guest)) {
        return p->Visit(visited);
    }
}

```

```

        }
        return ReturnType();
    }
};

#define DEFINE_VISITABLE() \
    virtual ReturnType Accept(BaseVisitor& guest) \
    { return AcceptImpl(*this, guest); }

```

- Using the library

```

class DocElement : public BaseVisitable<> {
public:
    DEFINE_VISITABLE()
};

class Paragraph : public DocElement {
public:
    DEFINE_VISITABLE()
};

```

```

class MyConcreteVisitor :
public BaseVisitor, // required
public Visitor<DocElement>, // visits DocElements
public Visitor<Paragraph> // visits Paragraphs {
    public:
        void Visit(DocElement&) {
            std::cout << "Visit(DocElement&) \n";
        }
        void Visit(Paragraph&) {
            std::cout << "Visit(Paragraph&) \n";
        }
};

```

```
int main() {  
    MyConcreteVisitor visitor;  
    Paragraph par;  
    DocElement* d = &par;  
                                // "hide" the static type of  
                                // 'par' d->Accept(visitor);  
}
```

This implementation is adequate for most situations. However, for applications where speed is a requirement, the **dynamic_cast** may be inhibitive.

And the Cyclic Visitor?

This is also referred to as the GoF Visitor

The bottom line is, we have a collection of types we want to visit. Therefore, pass a typelist to a Cyclic Visitor

```
// Forward declarations needed by the typelist  
class DocElement;  
class Paragraph;  
class RasterBitmap;  
  
// Visits DocElement, Paragraph and RasterBitmap  
typedef CycleVisitor <  
    void, // return type  
    TYPELIST_3(DocElement, Paragraph, RasterBitmap)  
> MyVisitor;
```

Follow a similar technique as was applied in Abstract Factory

```
template <typename R, class TList>
class CyclicVisitor : public GenScatterHierarchy<
    Tlist,
    Private::VisitorBinder<R>::Result> {
    typedef R ReturnType;
    template <class Visited>
    ReturnType Visit(Visited& host) {
        Visitor<Visisted>& subObj = *this;
        return subObj.Visit(host);
    }
};
```

- CyclicVisitor inherits a Visitor<T> for each type T in the typelist TList

The visitable hierarchy is set up by

```
typedef CycleVisitor <
    void, // return type
    TYPELIST_3(DocElement, Paragraph, RasterBitmap)
> MyVisitor;
```

```
class DocElement {
    public:
        virtual void Accept(MyVisitor&) = 0;
};
```

```
class Paragraph : public DocElement {
    public:
        DEFINE_CYCLIC_VISITABLE(MyVisitor);
};
```

Putting it all together?

Defining the catch-all function

- For the cyclic visitor, this is easy. In the typelist add the class `DocElement` at the head. This will result in `Visit(DocElement&)` to be defined which will be called when trying to visit an unknown type.
- With the acyclic visitor, instead of *visiting an unknown class by a known visitor*, a *unknown visitor visits a known class*. This is as a result of how the `BaseVisitable::AcceptImpl` is defined. It will return a default value of `Returntype()`. Require a policy to define the catch-all behaviour. `BaseVisitable` changes as follows:

```
template <
    typename R = void,
    template <typename, class>
        class CatchAll = DefaultCatchAll >
class BaseVisitable
```

```

{
... as above ...
template <class T>
static ReturnTpe AcceptImpl(T& visited,
    BaseVisitor& guest)  {
    if (Visitor<T>* p =
        dynamic_cast<Visitor<T>*>(&guest))
    {
        return p->Visit(visited);
    }
    // Changed
    return CatchAll<R, T>::OnUnknownVisitor(visited,
                                                guest);
} };

```

```
template <class R, class Visited>
struct DefaultCatchAll
{
    static R OnUnknownVisitor(Visited&, BaseVisitor&)
    {
        // Here's the action that will be taken when
        // there is a mismatch between a Visitor and
        // a Visited. The stock behaviour is to return
        // a default value
        return R();
    }
};
```